

Opening the Gate: Multiracial Identity and Institutional Access in Congressional Primaries*

Kaiyla Dartey[†] Isaiah C. Johnson[‡]

July 4, 2026

Abstract

Multiracial-identified Black candidates are often framed as uniquely positioned coalition builders, able to attract cross-racial support that their Monoracial peers cannot. We test this claim with an original dataset of Black candidates in U.S. House primaries from 2016 to 2024, the first to systematically distinguish Monoracial and Multiracial Black candidates at the congressional level. Using a novel LLM-assisted racial identification methodology validated against a human-coded gold standard, we find that Multiracial candidates raise more money, earn more institutional endorsements, and win primaries at higher rates. Yet this premium is sharply gendered. Multiracial men outperform all groups, winning roughly 67% of their primaries, while Multiracial women win just 33%, the lowest rate of any group. Monoracial women, by contrast, outperform Monoracial men. We explain this reversal with a theory of intersectional gatekeeping, in which donors, party organizations, and identity-based PACs read racial and gender signals together rather than separately. For men, Multiracial identification signals expanded electability. For women, it disrupts a Monoracial credentialing logic that Black-specific institutions consolidated around Black women after 2016. The findings show that the presumed advantages of Multiracial identification are neither universal nor unconditional, with implications for descriptive representation, electoral gatekeeping, and the measurement of racial identity.

*. The authors would like to thank the members of The Lab on Politics, Race and Experimental Methods. Both authors contributed equally to this manuscript.

†. Ph.D. Candidate, Princeton University. kdartey@princeton.edu.

‡. Ph.D. Candidate, Princeton University. ij17911@princeton.edu.

1 Introduction

Black representation in the United States (U.S.) Congress has expanded considerably over the past two decades, but expansion is not the same as incorporation. As the number of Black candidates competing in congressional primaries has grown, so too has the visibility of a particular candidate profile: Black politicians who publicly identify as Multiracial or of mixed racial heritage (Hardy-Fanta et al. 2013). From Barack Obama’s presidential campaigns to a new generation of congressional contenders, Multiracial-identified Black politicians have attracted considerable attention from scholars, journalists, and political strategists alike. Much of this attention centers on a compelling intuition, that Multiracial candidates are uniquely positioned to build cross-racial coalitions, attracting support from white voters and majority-white institutions in ways that their Monoracial Black counterparts may not (Lemi 2017, 2021, 2026; Velasquez-Manoff 2017). In this view, Multiracial identification is an electoral resource, a signal of broad appeal that translates into fundraising advantages, institutional endorsements, and ultimately electoral success.

We put the coalition-building hypothesis to a direct empirical test. Drawing on an original dataset of Black candidates competing in U.S. House primaries from 2016–2024, the first to systematically distinguish Monoracial and Multiracial Black candidates at the congressional level, we find that the story is considerably more complicated. The dataset, assembled using a novel artificial intelligence (AI) assisted racial identification methodology, spans nearly a decade of congressional elections. At first glance, the evidence appears to confirm the coalition-building story. Multiracial Black candidates raise dramatically more money than their Monoracial peers—median pre-primary receipts of roughly \$905,000 compared to \$114,000—and receive endorsements from both Black-specific institutional gatekeepers, such as the Congressional Black Caucus Political Action Committee (CBC PAC) and the Collective PAC, and general institutional funders, such as the Democratic Congressional Campaign Committee (DCCC), at markedly higher rates. They also win their primary elections more often. If one looked only at these aggregate patterns, the coalition-building hypothesis would look compelling.

But the aggregate pattern conceals a far more complicated story. When we examine the Multiracial premium through the lens of gender, the picture fractures. Among men, the premium is large and consistent. Multiracial Black men raise more money from more donors, earn endorsements at far higher rates, and win their primaries at roughly 67%, more than twenty percentage points above Monoracial Black men. Among women, however, the premium disappears and then inverts. Multiracial Black women win their primaries at just 33%, the lowest rate of any group in our data. Monoracial Black women, by contrast, win at over 51%, outperforming not only Multiracial Black women but Monoracial Black men as well. The signal that unlocks resources and institutional support for Multiracial men appears, for Multiracial women, to close doors rather than open them.

This reversal poses a puzzle that aggregate stories of Multiracial political incorporation cannot resolve. If Multiracial identification signals cross-racial electability, why should that

signal benefit men but harm women? Why would Black-specific institutions like the CBC PAC and Collective PAC, organizations one might expect to center Monoracial Black candidates, reward Multiracial men at dramatically higher rates while extending no equivalent premium to Multiracial women? And why do Monoracial Black women, whose political emergence has been among the most celebrated developments in recent electoral cycles (Brown 2014b; Brown, Clark, and Mahoney 2022), outperform their Multiracial counterparts by such a striking margin?

We argue the answer lies in what we term intersectional gatekeeping: the process by which institutional actors, specifically donors, party organizations, and identity-based PACs, evaluate candidate viability not through race or gender alone, but through their simultaneous intersection (Crenshaw 1989). Multiracial identification does not carry a fixed electoral value. Its meaning is constructed in context, interpreted differently depending on who is doing the evaluating and who is being evaluated. For Black men, Multiracial identification is legible to institutional gatekeepers as a signal of cross-racial coalition-building capacity, a form of expanded electability they are willing to invest in. For Black women, however, the same signal interacts with existing gender dynamics in the political arena in ways that complicate rather than amplify perceived viability. Black women have built a distinctive politics of identity-linked representation in which Black-specific institutions have come to treat Monoracial Black women as the credentialed representatives of Black communities, a credentialing logic that Multiracial Black women do not easily access. The result is that Multiracial identification produces sharply divergent outcomes depending on the gender of the candidate bearing it.

This paper also clarifies the temporal dynamics we observe in our data. The post-2016 period brought a dramatic reorientation of Black-specific institutional support toward Black women, a shift we trace to the dual shocks of the #MeToo movement and the historic surge in Black women’s candidacies during the 2018 midterm cycle (Simien 2022, 2023). Yet this reorientation accrued primarily to Monoracial Black women. Multiracial Black women were largely left outside of it, consistent with an account in which institutions mobilizing in support of Black women’s representation were doing so through a framework that centered Monoracial identity as the relevant political credential.

Our findings speak to several bodies of literature. We advance scholarship on Multiracial political incorporation by demonstrating that the presumed electoral advantages of Multiracial identification are neither universal nor unconditional (Lemi 2021). We contribute to work on Black women’s political emergence by revealing that aggregate gains in Black women’s representation mask important variation in which Black women attract institutional support and electoral access (Brown 2014b; Smooth 2021). And we contribute to the literature on electoral gatekeeping by showing how institutions process intersecting identity signals in ways that produce divergent, and sometimes perverse, outcomes for candidates whose identities complicate simple racial or gender categorizations (Crowder-Meyer 2013; Shah 2014).

The paper proceeds as follows. Section 2 develops our theoretical framework, drawing on literature from Multiracial politics, descriptive representation, and intersectionality to derive

expectations about the gendered dynamics of institutional support. Section 3 describes our data and the AI-assisted racial identification methodology. Section 4 discusses our validation strategy, introducing results of our concurrent and predictive validity tests. Section 5 presents our empirical findings, moving from descriptive patterns in candidate geography, fundraising, and institutional support to the gendered reversal at the center of our argument. Section 6 concludes by discussing the implications for racial representation, electoral gatekeeping, and the politics of Multiracial identity in contemporary American democracy.

2 Theory

Our theoretical argument proceeds in three steps. We begin with the prevailing account of Multiracial identity as an electoral resource, which frames Multiracial candidates as coalition builders and motivates institutional investment in them. We then turn to the distinctive position of Black women in electoral politics, tracing how Black-specific institutions have developed a credentialing logic that treats Monoracial Black women as the recognized representatives of Black communities. Finally, we integrate these two accounts into a theory of intersectional gatekeeping, which explains why the same signal that expands institutional access for Multiracial men forecloses it for mixed-race women. We close the section by describing the measurement challenge that has kept these dynamics from empirical view. We propose that our methodology, described in Section 3, is designed to solve this challenge.

2.1 Multiracial Identity as an Electoral Resource, and Its Limits

Multiracial candidates are unique in that they are framed as coalition builders due in part to their position straddling two (or more) identities (Lemi 2017, 2021, 2026; Velasquez-Manoff 2017). Multiracial candidates are also aware of this framing (Lemi 2017), strategically engaging with their identity both within institutions and with voters (Brown 2014a; Johnson 2026; Lemi 2026). And while there are qualitative accounts of Multiracial state legislators (Lemi 2017), we still know very little about this subgroup and the impact their rise has on political representation at the federal level.

One area where race and identity are understudied is within campaign finance (Grumbach and Sahn 2020). We know that well-financed candidates perform better electorally (Bonica 2017; Thomsen 2023). Interest groups and PACs play a significant role in electoral outcomes (Magee 2002). Moreover, PACs are strategic institutional players, supporting candidates who, for example, moderate their issue positions in order to improve electability (Meisels 2026; Myers 2025). Identity-based PACs, such as the CBC PAC, additionally serve as gatekeepers, socially monitoring group members and sanctioning them when they deviate (Rendleman 2023). Because Multiracial candidates are situated as uniquely able to navigate multiple racial domains, PACs may be more willing to strategically donate to them in order to improve their electoral odds. Importantly, identity-based PACs may also increase donations to Multiracials

in hopes that their social monitoring capabilities will rein in deviations, a concern rooted in fear that a mixed-race candidate may not fully align with the group (Lemi 2017). Because of the dearth of candidate-based information on Multiracials, these questions and downstream policy implications have gone unanswered.

Yet the coalition-building account, as it currently stands, treats Multiracial identification among candidates as a homogeneous signal, one conferring equal access to opportunity among men and women. We argue that this assumption is wrong. The value of the Multiracial signal is constructed in context, and one of the most consequential contexts in electoral politics is gender (Brown 2014a; Lemi and Brown 2019). To see why, we must first examine the distinctive institutional position that Black women occupy in American electoral politics.

2.2 Black Women in Electoral Politics: Compound Disadvantage, Representational Credentialing, and the Infrastructure of Institutional Support

The study of Black women’s political incorporation begins from a well-documented paradox: Black women are among the most politically active and consistently Democratic constituencies in the American electorate, a cohesion rooted in the shared social experience of navigating racial marginalization and enforced through community-level norms of political solidarity, yet they have historically faced some of the steepest barriers to office-holding of any demographic group (Brown 2014b; Dawson 1994; Smooth 2021; White and Laird 2020). This paradox reflects the structural position Black women occupy in American politics, one defined not simply by the compounding of racial and gender disadvantages, but by the specific and irreducible form of marginalization produced when those two axes of inequality intersect. The theoretical significance of this interaction was established by Crenshaw’s (1989) foundational account of intersectionality, which demonstrated that the social positions created by the convergence of race and gender cannot be recovered by treating either dimension independently. Black women’s political exclusion is not equivalent to what Black men face plus what white women face. It is a qualitatively distinct condition produced by the intersection of racial and gendered expectations in ways that neither race nor gender scholarship alone can adequately characterize (Collins 2000; Reingold, Haynie, and Widner 2020).

Those barriers to electoral participation and office-holding are substantial and well-documented. Black women candidates must navigate a candidate evaluation environment in which the default image of a credible congressional candidate has historically been constructed as simultaneously white and male, an image embedded in the heuristics through which parties recruit candidates, donors evaluate viability, and media cover campaigns (Crowder-Meyer 2013; Sanbonmatsu 2006). Racial minority candidates face lower baseline perceptions of electability in majority-white districts and encounter donor networks less practiced at evaluating their viability, producing systematic fundraising disadvantages that compound at the primary stage before a general election audience is even reached (Philpot and Walton Jr 2007; Shah 2014). Women candidates simultaneously face gendered fundraising constraints, differential and often

discouraging treatment by party recruiters operating from assumptions about which candidate profiles are electorally viable, and persistent stereotypes about issue competence and executive leadership capacity that shape both voter evaluations and gatekeeper investment decisions (Crowder-Meyer 2013; Dittmar, Sanbonmatsu, and Carroll 2018; Sanbonmatsu 2002). For Black women, these disadvantages do not simply add together; they interact in ways that can compound the burden of establishing credibility across both dimensions simultaneously. A Black woman seeking institutional support must convince gatekeepers of her viability as a Black candidate and as a woman candidate, within institutional frameworks that have historically under-invested in both, and that were not built with either in mind.

Yet Black women have emerged over the past two decades as an increasingly prominent force in congressional politics, a development that requires explanation beyond the simple lifting of structural barriers. We argue, alongside an emerging literature on Black women’s political emergence, that this development is partly the product of a distinctive form of representational credentialing that has developed within and around Black political institutions. Drawing on Mansbridge’s (1999) contingent theory of descriptive representation, which holds that disadvantaged groups gain the greatest representational value from members of their own group in conditions of group mistrust, uncrystallized interests, and contested political legitimacy, we observe that Black women’s intersectional identity generates a specific and institutionally recognizable claim to representational authority. As Brown (2014b) documents in her foundational account of Black women legislators’ self-understanding, the intersection of race and gender does not function only as compound disadvantage: it also generates a political standpoint, rooted in the experience of navigating both racial and gender marginalization, from which a distinctive form of representational capacity is constructed and understood (Collins 2000). Black women politicians have built a political landscape in which their intersectional identity is itself constitutive of a particular claim to represent. Black women are not merely representatives of Black communities who happen to be women, or representatives of women who happen to be Black; they occupy a representational position whose authority derives precisely from the intersection, a position that Dawson’s (1994) framework of linked fate helps illuminate by showing how the shared structural experience of racial subordination generates a basis for political solidarity and group-based accountability.

This credentialing logic has become progressively institutionalized through the emergence and scaling of Black-specific political organizations. The CBC PAC; the Collective PAC, established in 2016 explicitly to build infrastructure for electing Black candidates at every level of government; and Higher Heights for America, dedicated specifically to Black women’s political leadership, have together created an institutional ecosystem through which the credentialing of Black women as candidates with particular representational authority is translated into material resources: fundraising, endorsements, strategic training, and earned media amplification. These organizations do not invest indiscriminately in Black women candidates; they invest through evaluative frameworks in which racial identity, demonstrated community accountability, and representational commitment function as implicit criteria of viability alongside more conventional measures of electoral competitiveness (Brown, Clark, and Mahoney 2022; Slaughter, Crowder, and Greer 2024). The effect is to consolidate and institutionalize a credentialing logic that preferentially rewards candidates whose identity

signals unambiguous membership in and accountability to the Black communities they seek to represent.

The post-2016 political environment sharpened and formalized this logic considerably. The mobilization generated by the #MeToo movement intersected with a historic surge in Black women’s candidacies during the 2018 midterm cycle to produce a moment of symbolic and institutional reorientation, one in which Black-specific organizations substantially increased the scale and intentionality of their investment in Black women candidates, and in which the representational credential embedded in Black women’s intersectional identity became an increasingly explicit rather than implicit criterion of organizational support (Simien 2022, 2023; Smooth 2021). The result was to consolidate a credentialing architecture within Black political institutions, one in which Monoracial Black identity had become, if not the explicit criterion, a central implicit organizing logic of institutional investment in Black women specifically.

The theoretical implication that follows from this account, and which the existing literature has not yet examined, is that the credentialing logic embedded in Black-specific institutional gatekeeping creates a structured expectation about what Black women candidates look like—an expectation that Multiracial Black women may satisfy only partially, if at all. If Black-specific institutions are investing in Black women candidates in part because their intersectional identity constitutes a distinctive representational credential grounded in Monoracial Black experience and gender marginalization, then Multiracial Black women present a complication the current literature has not yet theorized. Multiracial women are gendered in the way that institutions have learned to invest in, but racially identified in a way that disrupts the specific Monoracial credential driving that investment. To understand how institutions navigate this complication, and why it produces such sharply divergent electoral outcomes, we need a theoretical framework capable of modeling how multiple identity signals interact in the evaluative logic of electoral gatekeepers. We develop that framework in the following section.

2.3 Intersectional Gatekeeping: A Theory of How Institutions Read Race and Gender Together

The concept of intersectional gatekeeping begins from Crenshaw’s (1989) core insight but applies it to an explicitly institutional rather than purely social context. Where Crenshaw demonstrated that legal and social systems governing anti-discrimination treated race and gender as independent axes, and thereby systematically obscured the distinctive position of Black women who experienced both, we argue that political institutions and electoral gatekeepers engage in an analogous, consequential error of independent processing. Donors, party organizations, identity-based political organizations, and the broader ecosystem of electoral gatekeeping do not evaluate candidates’ racial and gender identities separately and then sum the implications; they read these identities together, producing evaluations that are qualitatively shaped by their intersection in ways that cannot be derived from the analysis of

either dimension alone (Reingold, Haynie, and Widner 2020). The signal value of Multiracial identification for a Black candidate is not a stable quantity that is simply appended to whatever signal gender produces. It is conditioned by gender in ways that can alter not merely the magnitude but the direction of the effect.

This relational quality of identity signal processing is the foundation of our argument. When institutional gatekeepers evaluate a candidate’s viability, they draw on heuristics built from prior experience with which kinds of candidates in which kinds of districts attract which kinds of support and win; heuristics that are neither racially nor gender neutral but that encode prior patterns of investment and success in ways that embed existing assumptions about electability (Crowder-Meyer 2013; Sanbonmatsu 2006; Shah 2014). Multiracial identification enters this evaluative framework not as an abstract attribute with a fixed value but as a contextually interpreted signal whose implications are read through the lens of other visible candidate characteristics, most centrally, gender. The same racial identity signal can therefore produce systematically different evaluative responses depending on the gender of the candidate bearing it, because the institutional frameworks through which that signal is interpreted are themselves gendered.

For Multiracial Black men, we argue that racial and gender signals interact in a mutually reinforcing and institutionally legible way. Multiracial identification among Black male candidates is interpretable to institutional gatekeepers, including Black-specific ones, as a signal of expanded coalition capacity: the demonstrated or presumed ability to attract cross-racial support, compete viably in majority-white or racially diverse districts, and build donor coalitions that extend beyond majority-Black networks (Lemi 2017, 2021, 2026; Velasquez-Manoff 2017). This reading of Multiracial identification as coalition-broadening asset is available to gatekeepers because it is not complicated by gender dynamics. Male candidates remain the institutional default in congressional elections, the candidate type around which most gatekeeping heuristics were originally constructed, and male gender therefore poses no independent credentialing challenge that would complicate the reading of Multiracial identification as simply an electability asset. Black-specific gatekeepers face a further, distinct consideration, supporting Multiracial male candidates involves accepting some strategic risk that the candidate may emphasize cross-racial coalition building over the robust substantive representation of Black community interests specifically, the core concern animating linked fate–driven Black political investment (Dawson 1994; Mansbridge 1999). But this risk is understood institutionally as a strategic trade-off rather than an identity disqualification. The candidate’s gender does not disrupt the evaluative framework through which their racial identity is being processed. The Multiracial premium for Black men is the product of this specific institutional configuration. As a result, gender creates no credentialing obstacle, and Multiracial identification is read as expanding electability. Therefore, institutional gatekeepers calculate the strategic trade-off as worth making.

For Multiracial Black women, the dynamics operate in a fundamentally different register, and the theory predicts not simply a reduced premium but a reversal. The credentialing logic that has come to drive Black-specific institutional investment in Black women is explicitly intersectional in its structure. The credential rewards candidates whose identity signals

accountability to Black communities grounded in the shared experience of Monoracial Black identity and gender marginalization simultaneously (Brown 2014b; Collins 2000; Mansbridge 1999). Multiracial Black women are positioned awkwardly within this framework in a way that Multiracial Black men are not. Their gender marks them as candidates Black-specific institutions have built infrastructure to support; their Multiracial identification disrupts the racial dimension of the credential that drives that support. This is not a case where one identity advantage partially offsets another identity disadvantage, producing a neutral or attenuated net effect. It is a case where the two signals interact to produce a specific form of institutional illegibility. The Multiracial Black woman candidate does not fit cleanly into the representational framework through which Black-specific institutions have learned to identify and invest in Black women, and she does not fit the coalition-building framework through which Multiracial men’s expanded electability is rewarded either. The result is that gatekeepers operating from heuristics built for a different candidate profile respond with reduced investment, producing the electoral reversal at the center of our empirical findings.

The institutional specificity of this dynamic matters theoretically and generates a testable prediction about where the reversal should be most pronounced. Black-specific gatekeepers and general institutional funders process identity signals through different evaluative frameworks, and intersectional gatekeeping theory predicts that the gender asymmetry in the Multiracial premium will be most acute among Black-specific institutions precisely because their investment rationale is most explicitly anchored in the credentialing logic that Multiracial identification disrupts. The CBC PAC, Collective PAC, and Higher Heights operate through an explicit institutional commitment to Black representational interests, making the Monoracial credential most central to their evaluative logic (Brown, Clark, and Mahoney 2022; Slaughter, Crowder, and Greer 2024). General institutional funders such as the DCCC and organizations like EMILY’s List operate through a somewhat different framework centered on perceived electoral viability rather than representational credentialing per se, but they are not immune from analogous dynamics (Dittmar, Sanbonmatsu, and Carroll 2018; Sanbonmatsu 2006). For these organizations, Multiracial identification in Black men may signal cross-racial viability in competitive districts; for Black women, the combination of gender and Multiracial identification may produce evaluative ambiguity, as the candidate does not map cleanly onto the identity templates through which electoral viability in majority-Black or competitive districts has been learned and rewarded. The asymmetry should therefore be present across both gatekeeper types, but most theoretically acute and empirically pronounced among Black-specific institutions whose investment logic is most explicitly identity-linked.

Finally, intersectional gatekeeping theory generates a temporal prediction that extends the analysis beyond cross-sectional patterns. If the consolidation of Black-specific institutional investment in Monoracial Black women accelerated after 2016, as the account in Section 2.2 argues, then the gap between Monoracial and Multiracial Black women in institutional support and electoral outcomes should widen in post-2018 election cycles relative to earlier ones. The post-2018 reorientation of Black-specific gatekeeping toward women, driven by the mobilization dynamics of the #MeToo moment and the 2018 surge in Black women’s candidacies (Simien 2022, 2023; Smooth 2021), should have disproportionately benefited Monoracial Black women, whose identity aligned with the credential being institutionally consolidated,

relative to Multiracial Black women, deepening the latter’s institutional disadvantage over time. This temporal prediction is directly testable in the data we have assembled, and its confirmation would provide additional evidence that the mechanism generating the reversal is the institutionalization of a Monoracial credentialing logic, rather than some static property of voter or gatekeeper preferences that predates the period we study.

2.4 From Theory to Measurement

Testing these predictions requires data that no existing study provides, a systematic, individual-level accounting of which Black congressional candidates identify as Multiracial and which as Monoracial. The obstacle is, at bottom, a measurement problem (Denton and Deane 2010; Sen and Wasow 2016; Snipp 2003).¹ While a set of scholars utilize candidates to study political phenomena with a binary classification of race (Hassell and Visalvanich 2019; Porter and Treul 2025), several scholars have opted for a hand-coded classification approach to the racial identity of political candidates. At the federal level, Fraga (2016) and Fraga and Hassell (2021) develop a coding scheme for racial identity via gathering information from news articles, candidate statements and websites, membership in identity-based political institutions, among other resources. Similarly, Grumbach and Sahn (2020) extend Goggin (2017) and code racial identity from 1980–2012 to study racial inequality in campaign finance. In early work examining the role of race in the primary process, Branton (2009) compiles information on candidates’ racial background from several resources. A larger body of work studying racial identity of politicians at the state level also utilizes aggregate information to hand-code a measure of race (Atsusaka 2021; Fraga, Gonzalez Juenke, and Shah 2020; Fraga, Shah, and Juenke 2020; Juenke 2014; Juenke and Shah 2016; Lublin and Wright 2024; Shah, Scott, and Juenke 2019). Yet, this research agenda largely ignores the role of Multiracial candidates (Lemi 2017), severely limiting our knowledge of this fast-growing subpopulation (Hardy-Fanta et al. 2013).

A recent alternative bypasses hand coding entirely. Dasanaik (2026) utilizes the ability of large language models (LLMs) to conduct name-pattern recognition for their racial identification measure, effectively an LLM-based analogue to the Bayesian surname approach of Imai and Khanna (2016) and Imai, Olivella, and Rosenman (2022). However, the underlying measurement logic infers race probabilistically from a proxy rather than documenting it at the individual level. Race, like gender, is a construct (Omi and Winant 1986), affected by one’s self-identification and socialization (Davenport 2016; Thomas and Speight 1999). Such a construct leaves its evidence in a person’s biography, not their name alone.

Indeed, centuries of racialization in the U.S. have made it difficult to ascertain *who* is

1. We agree with the critique of Multiracialism not as a measurement issue needing to be solved but as a continued attempt to racialize and quantify a constructed concept (Fields and Fields 2012). Nonetheless, to be agnostic toward identification, particularly among candidates who *choose* to delineate themselves from the Monoracial group, does a disservice to answering questions rooted in democratic theory. The presumption that “all skin-folk are kin-folk” ignores nuances that disproportionately affect representation (Dovi 2002; Wamble 2025).

Black from surname alone (Lockhart, King, and Munsch 2023). This is the exact reason for the high false negative rates found in Imai and Khanna (2016), and why their model was improved to account for measurement error (Imai, Olivella, and Rosenman 2022). And while adding covariates for geography improves model performance (Imai and Khanna 2016; Imai, Olivella, and Rosenman 2022), these models still introduce bias for minority subpopulations (Greengard and Gelman 2026). Misclassification tracks socioeconomic status, such that minority respondents in predominantly White areas are routinely mislabeled as White (Argyle and Barber 2024). These failures are especially consequential for the subpopulation we study. For Multiracial Americans, name-based inference is not merely noisy but conceptually inapplicable, as a surname reflects at most one branch of an identity formed through self-identification and socialization.² We opt against name-based inference altogether, replacing probabilistic-based models with individual-level documentation, classifying each candidate on direct biographical evidence of ancestry and self-identification, including candidates’ own statements, documented parentage, and contemporaneous reporting. In the next section, we provide an overview of this unique methodology, discussing the implementation process, including the coding pipeline, the development of a documentation tier system, and the construction of our racial identity taxonomy.

3 Methodology

Our approach utilizes LLMs to measure racial identity as a codebook (Halterman and Keith 2026), augmented for information acquisition. Following the guidance for incorporating human-coded “gold standard” validation when working with LLMs and machine learning techniques more broadly (Grimmer and Stewart 2013), we developed a rule-based codebook of identity classification that was first given to undergraduate and graduate research assistants (RAs) serving as our gold standard. These human coders provided a first screen of the 1,213 African American candidates found between 2016 and 2024, coding their racial background, gender, and acquiring other critical variables.³ The RAs followed a similar pattern as other human-based coding algorithms, utilizing news articles, campaign statements, and other sources commonly found through search engines to locate the racial identity of candidates. Given that recent developments in LLMs have led to models outperforming human-based or expert coding (Benoit et al. 2025; Gilardi, Alizadeh, and Kubli 2023; Heseltine and Clemm von Hohenberg 2024; Törnberg 2023), we sought to test the performance of generative AI as an alternative measurement tool for racial identity.

Because internet browser and software-based LLMs are unable to execute thousands of standardized, independently parameterized searches under a fixed instruction set (Benoit et al. 2025), we chose to utilize Anthropic’s application programming interface (API) to

2. Imai and Khanna (2016) and Imai, Olivella, and Rosenman (2022) and Dasanaik (2026) all relegate Multiracial observations to an “other” category. The high false negative rate within this category biases any downstream analyses.

3. For the purposes of this study, only African American-identified candidates are utilized. However, these data are part of a larger project coding the racial identity of all candidates from 1971–2024.

conduct a series of batch searches for candidate racial identifiers. In order to properly communicate with precision the steps required by the LLM to code for race, we developed a Claude skill that incorporated four rules:

1. Classification rests exclusively on explicit self-identification, documented parentage, or grandparent-level ancestry the candidate acknowledges.
2. Every candidate is recoded into newly created variables, leaving any preexisting racial data untouched.
3. Every classification receives a documentation-quality tier, recording how well-documented the coding is independently of what the coding is.
4. When searches cannot establish ancestry, the candidate is coded indeterminate rather than force-classified.

The first rule stems from canonical work on Multiracial identification, which utilizes both ancestral and self-identification measures of Multiracialism (Davenport 2016; Leslie 2022; Rockquemore and Brunsma 2008). The creation of a new racial identity variable ensured that there was no overwriting of the human-coded measure; moreover, this helped create a distinct classification metric used in validation techniques. The tier system, visualized in Table 1, allowed the model to develop a confidence scheme to help lower the false positive rate. Candidates were labeled as high confidence when identity was explicit, medium if evidence came from a single source or classification was not explicit, and low if there was no usable racial information. Finally, the LLM was instructed to label tier 3 candidates as indeterminate (NA), again to guard against false positives. In effect, the skill functions as a standing instruction set that governs every request the LLM receives, ensuring each candidate is coded under identical rules rather than ad hoc prompting.


Table 1: Documentation-Quality Tier System

Tier	Evidentiary standard	Confidence	Disposition
1	Ancestry explicitly documented in reliable sources: both parents documented, or authoritative self-identification.	HIGH	Coded; no further action
2	Thinner or single-source evidence; classification plausible.	MEDIUM/LOW	Coded; flagged for validation sample
3	No usable ancestry information located after searching.	LOW	Coded <i>indeterminate</i> ; routed to broader-search retry

Note: Tiers record how well-documented a classification is, independent of the classification itself. Tiers are assigned mechanically from the coding verdict: Tier 3 if indeterminate or no sources, Tier 1 if confidence is HIGH, Tier 2 otherwise. Tier 2 rows use MEDIUM confidence, or LOW when supported by at least one real source.

To identify the race of candidates, we instructed the LLM via the skill file to follow a taxonomy of racial classification. Figure 1 describes the taxonomy as a flowchart. First, the model searches for evidence of a candidate’s racial identity. Following the significance of self-identification in research on Multiracial politicians (Lemi 2017, 2026), explicit mention from the individual that they are mixed-race—whether through one’s biography, interview, or other form of communication—was the first step in the taxonomy. All candidates who self-identified as mixed-race were labeled as Multiracial. If self-identification was not present, the model searched for ancestral identification via parentage or grandparentage (Davenport 2016; Leslie 2022).⁴ If there was documented mixed-race status via ancestral means, then this candidate was recorded as Multiracial. If self-identification or documentation of ancestry determined they had two Monoracial parents, they were then coded as Monoracial Black. Otherwise, the model was instructed to return an indeterminate label.

Figure 1: Taxonomy of Identification



Note: Classification proceeds along two evidentiary routes: explicit self-identification, which classifies a candidate as Multiracial directly, and documented ancestral identification through parentage or acknowledged grandparentage.

All batch searches took place in June of 2026 utilizing Claude Opus 4.8, a proprietary

4. We acknowledge the debate regarding selective hypodescent (Spencer 2004). Nonetheless, we attempted to balance inclusion of multi-generational identities while circumscribing the broad claim that everyone in America is of some mixed-race ancestry (Fields and Fields 2012).

closed-source model. Table 2 provides an overview of the LLM coding pipeline. An initial pilot of 25 candidates was conducted to test the skill file, ensure there were no coding errors, and check the validity of a small sample of LLM-coded candidates. A full run was conducted on the remaining 614 candidates whose race was not previously known.⁵ After the initial LLM pass, the remaining 393 candidates who returned no identifiable racial information in the first round were retried a second time.⁶ The skill instructed Claude to conduct a maximum of five searches per candidate, though the realized average was roughly half a search per candidate. In total, 305 searches were conducted in roughly an hour of total usage time, at a total cost of \$14.82.⁷

Table 2: LLM Coding Pipeline: Usage and Cost

	Pilot	Main pass + retry	Total
Requests submitted	25	1,007	1,032
Main pass	–	614	614
Tier-3 retry	–	393	393
Web searches	9	296	305
Input tokens	153,776	6,105,881	6,259,657
Output tokens	3,251	148,039	151,290
Cost: model tokens (\$)	0.34	13.36	13.70
Cost: web search (\$)	0.04	1.08	1.12
Cost: total (\$)	0.38	14.44	14.82
Searches per candidate			0.50
Cost per candidate (\$)			0.024

Note: All requests used Claude Opus 4.8 (Anthropic) via the Batch API with the web search tool enabled, June 2026. Costs are billed API charges at batch rates.

4 Validating LLM-Based Racial Identification

Before providing validity checks for our LLM-based method, it is helpful to get an overview of where the LLM searches for a candidate’s racial identity. Within the Claude skill, every search was required to record the source from which it obtained information about the candidate’s racial identity. Table 3 shows the sources categorized into common themes. The first column covers all LLM-coded candidates in our dataset, with the second and third columns further delineating between Multiracial- and Monoracial-identified candidates.

5. Of the 1,213 candidates, roughly one-third were well-known incumbents whose race is documented. As such, the remaining 614 were predominantly candidates with no prior history of congressional office at the federal level.

6. The bulk of these candidates were tier 3. As a result, only 168 candidates were identified successfully in the second round of searches.

7. Rounded, this comes out to about \$15 an hour for approximately 300 searches. In contrast, a human coder at the same hourly rate is unlikely to accomplish thorough searches at this magnitude.

Roughly 79% of all sources were found in reference encyclopedias such as Wikipedia and Ballotpedia. The remaining categories rounding out the top five include news media (49%), government sources such as House.gov (25%), campaign materials (25%), and voter guides (9%). Interestingly, while the Multiracial and Monoracial source profiles are largely similar, two distinct channels appear: government sources and campaign materials. Multiracials are less likely to be identified in government sources than Monoracial candidates (19% versus 26%). Conversely, mixed-race candidates are more likely to be found by the LLM via their campaign materials (49% versus 23%). This provides early evidence that identifying as Multiracial is a selection process that mixed-race candidates engage in strategically—it takes an active choice to identify as Multiracial, and campaign materials are the most public way to delineate themselves to their constituency.

Table 3: Evidentiary Sources Behind the LLM Coding, by Category

Source category	All coded	Multiracial calls	Monoracial calls
Reference encyclopedias (Wikipedia, Ballotpedia, BlackPast)	78.2%	64.7%	79.2%
News media	49.4%	47.1%	49.6%
Official government sources (House.gov, FEC, state)	25.3%	19.1%	25.8%
Campaign materials (own website, campaign bio)	24.6%	48.5%	22.8%
Voter guides & candidate directories	8.5%	1.5%	9.0%
Advocacy & organizational profiles	4.1%	2.9%	4.2%
Social & professional profiles (LinkedIn etc.)	3.4%	1.5%	3.6%
Books & author biographies	2.8%	1.5%	2.9%
Propagated from a prior candidate-cycle coding	1.9%	0.0%	2.1%
Limited or no sources found	2.3%	0.0%	2.5%

Note: Shares of LLM-coded candidate-cycles whose returned source list cites at least one source in each category (multi-label, so columns exceed 100%; mean 2.0 categories per candidate; 1.3% of source strings match no category). Categories assigned by keyword matching on the free-text source field returned with each coding. Columns 2 and 3 split candidate-cycles by the LLM’s own call (Multiracial vs. Monoracial).

4.1 Concurrent Validity

Next, concurrent validity examines the LLM coding scheme against the RA-coded gold standard. To measure concurrent validity, we examine precision, or the positive predictive value, which tells us how much we should believe a positive LLM label. Next, we examine recall, which measures how sensitive the model is. Specificity measures the true rejection rate for Monoracial candidates. Finally, accuracy is the overall agreement between the LLM and the gold standard. Confidence intervals (CIs) are exact Clopper–Pearson (Clopper and Pearson 1934). Figure 2 shows all four classification metrics, together with the F1 score discussed below. Recall is very high (0.97, 95% CI [0.88, 0.99]), indicating that the LLM catches nearly all true Multiracial candidates, missing only 2 of the 57. Though slightly more modest, precision is also relatively robust (0.81, 95% CI [0.70, 0.89]). The LLM coded 68 Multiracial candidates, of which 55 were confirmed by the human-coded classification,

resulting in 13 false positives.⁸ Specificity is very high, as expected. The pattern here is that the LLM slightly errs toward over-including Multiracial-identified candidates.⁹ While the model hardly ever misses, it does overstate compared to the gold standard.¹⁰




Figure 2: Concurrent validity of the LLM coding: classification metrics and F1 with 90% (thick) and 95% (thin) confidence intervals. *Note:* Intervals are exact Clopper–Pearson for the proportion-based metrics and percentile bootstraps over candidate-cycles for F1. Precision is the share of LLM Multiracial calls confirmed by the RA gold standard; recall is the share of gold-standard Multiracial candidate-cycles the LLM recovers; F1 is their harmonic mean.

Raw accuracy is extremely high; the LLM shows nearly perfect agreement with the human-coded racial identity labels. This is because Multiracials are a small share of the dataset. Chance alone would correctly classify the large Monoracial majority. As a result, Table A3 in Appendix Section B provides two agreement estimates that take the imbalance between groups into account: Cohen’s κ (Cohen 1960) and Gwet’s AC1 (Gwet 2008).¹¹ First, Cohen’s κ equals 0.87, with bootstrapped confidence intervals. The reason for the drop in agreement is that κ is depressed under skewed distributions (Feinstein and Cicchetti 1990; Gwet 2008).¹² Gwet’s AC1, which is designed to be robust to skewed prevalence (Gwet 2008),

8. Due to restricting the sample of human-coded labels to the subset that the LLM actually coded (eliminating 225 candidates), there are 12 gold-standard-labeled Multiracial candidates not included in the concurrent validity measures. However, this missingness is concentrated in tier-3 candidates—the lowest rung of confidence; therefore, Figure 2 provides a robust comparison between the LLM and human-coded labels.

9. A McNemar test formalizes this asymmetry in Table A3. With 13 false positives and 2 false negatives, the McNemar test results in a statistically significant imbalance ($p < 0.01$). This allows us to conclude that the error rate is not noise but directional: the LLM over-identifies Multiracial candidates.

10. Because of the 225 eliminated candidates, recall is conditional on the observations the LLM actually labeled. Unconditional recall is lower: $\frac{55}{(57+12)} = 0.80$.

11. See also McLaren et al. (2026) for the utilization of Cohen’s κ for binary LLM outcomes.

12. This phenomenon is known as the prevalence paradox. In particular, when the prevalence of one group

equals 0.98, indicating near-perfect agreement between the two coding schemes. While there is debate over the appropriate agreement estimate, both measures indicate that the LLM is an accurate replacement for human labels.

Finally, F1 captures the balance between precision and recall on the Multiracial-coded labels. F1 scores vary largely by the task and model types in the social sciences (Halterman and Keith 2026; Pangakis, Wolken, and Fasching 2023); however, high-performing zero-shot models often score above 0.7 (Ziems et al. 2024). Therefore, we set 0.7 as our benchmark. Results indicate that we well exceed this value, with tight bootstrapped confidence intervals (0.88, 95% CI [0.81, 0.94]). Therefore, overall model performance is strong, and the main takeaway is that LLMs have strong concurrent validity against gold-standard racial identity measures.

4.2 Predictive Validity

To assess predictive validity, we examine whether the LLM-based coding scheme reproduces the same downstream relationships as the gold-standard labels across a battery of outcomes associated with candidates’ racial identity. An additional set of campaign finance variables is included as well, since these appear in the downstream analyses of our substantive case study. Figure 3 reports both Pearson’s r and Spearman’s ρ , each computed across the full set of 11 outcome gaps; continuous outcomes are plotted as circles and binary outcomes as triangles.¹³ The y-axis is the standardized Multiracial–Monoracial gap under the LLM coding scheme, while the x-axis is the gap under the human-coded labels. The dashed 45-degree line represents agreement between the LLM and the human-based labels on a given outcome. Overall, there are high correlations between the LLM and gold-standard coding schemes. Both the Pearson and Spearman correlations are greater than 0.90 (bootstrapped confidence intervals), indicating strong linear and rank agreement. Pearson’s r is modestly attenuated relative to Spearman’s ρ , a gap driven by the continuous campaign finance measures: because these outcomes are heavily skewed, small differences between the two codings produce outsized deviations in the standardized gaps, penalizing the linear correlation more than the rank-based one.

Across both concurrent and predictive validity metrics, it is clear that the LLM-based racial identity labels approximate the human-based gold standard. The findings are equivalent to or exceed other LLM-to-human comparisons (Gilardi, Alizadeh, and Kubli 2023; Halterman and Keith 2026; Pangakis, Wolken, and Fasching 2023; Ziems et al. 2024). While we are cautious given that our LLM choice is closed-source with a zero-shot learning paradigm

is small (here, Multiracials), the function of κ becomes highly sensitive to the raters’ marginal rates and is penalized accordingly. κ is maximized when prevalence is balanced, so as prevalence moves toward either extreme the coefficient declines even when observed agreement is high.

13. Given the small number of variables analyzed, wide bootstrapped confidence intervals are reported in Appendix Table A4. See Hauke and Kossowski (2011) for a justification for reporting both Pearson and Spearman. See De Winter, Gosling, and Potter (2016) for a comparison of Pearson and Spearman pertaining to sample size adjustments.




Figure 3: Downstream-gap agreement under the two codings. *Note:* Each point is one outcome from the battery; the horizontal axis is the standardized Multiracial-minus-Monoracial gap under the RA gold standard, the vertical axis the same gap under the LLM coding (continuous outcomes: median gap scaled by the pooled interquartile range (IQR), circles; binary outcomes: percentage-point share gaps, triangles). The dashed line is the 45-degree line. Pearson and Spearman correlations with bootstrap 95% CIs over candidate-cycles are annotated.

(Halterman and Keith 2026; McLaren et al. 2026; Ollion et al. 2024; Spirling 2023), we opt to utilize the LLM-based coding scheme for our substantive case study that follows.¹⁴

14. Though, see Törnberg (2023) and Ziems et al. (2024) on zero-shot models.

5 Results

5.1 A Substantive Test of Intersectional Gatekeeping

Our methodology for coding Multiracial identification lets us examine where these candidates run, their partisan and ideological breakdown, and how institutional resources are allocated across identity groups. Together, these descriptive patterns motivate our central argument that Multiracial identification is not random. Identification is structured by geography and district context, and institutional resources are distributed in ways consistent with funding sources treating identity as informative—a premium that is, moreover, sharply gendered. While we utilize the LLM-based racial identity labels for the main analyses, Appendix Section B replicates all results with the gold-standard coding scheme. Findings are consistent across both identity measures.

Table 4 provides an overview of the data. Of the 1,213 candidates in the dataset, 988 were successfully coded using LLM batch searches (corresponding to 540 unique candidates, reflecting repeat candidacies across cycles). We identify 68 of these candidates as Multiracial and the remaining 920 as Monoracial. Across gender lines, roughly two-thirds of both groups are men. Both groups identify overwhelmingly as Democrats, though Multiracial candidates do so at a somewhat lower rate (71% versus 76%).

Table 4: Sample Overview: Black Congressional Candidate-Cycles, 2016–2024 (LLM Coding)

Measure	Multiracial	Monoracial	Total
All candidate-cycles	68	920	988
Share of coded sample	6.9%	93.1%	100.0%
Unique candidates			540
Men	44	545	589
Women	24	375	399
% Democrat	70.6%	75.9%	

Note: Counts of candidate-cycles and unique candidates, the group shares of the coded sample, the gender split, and the share identifying as Democrats, under the LLM coding on the head-to-head sample.

Yet their ideological profiles are relatively similar. Figure 4 shows the density of campaign finance (CF) scores (candidates) and DW-NOMINATE (incumbents) on a liberal to conservative dimension. Both groups are similarly liberal, though there are moderate bumps in conservatism for Multiracial-identified candidates (see Appendix Table A1 for the exact ideological placement by group).

Where Black candidates identify as Multiracial varies sharply by region. Figure 5 maps the Multiracial share of all Black-identified candidates by state. The pattern is strikingly sectional. The Multiracial-identified share is concentrated in the coastal Northeast and




Figure 4: Candidate ideal points by identity group: kernel densities of CF scores (all candidates) and first-dimension DW-NOMINATE (incumbents), restricted to $[-1, 1]$. Positive values are more conservative.

Mid-Atlantic, including New York, Virginia, Maryland, and Massachusetts, together with California and Washington; the south, on the other hand, remains largely Monoracial, following the legacy of hypodescent (Davis 1991; Leslie and Sears 2022). The contrast is not simply about where Black candidates run, but about how they identify conditional on running. This sectional divide echoes the historical geography of the one-drop rule, whose binary racial classification was most rigidly enforced in the Jim Crow South (Davis 1991; Leslie and Sears 2022). We read this as suggestive of a geographic logic to identification. Whether a Black candidate identifies as Multiracial appears conditioned by the racial context in which they run, not merely by underlying heritage.

This logic extends from geographic region to the racial context of individual districts. Table 5 depicts district composition, capturing the racial demographics of the voting age population (VAP) and the Democratic presidential vote share, along with whether the district is majority-Black or majority-minority. While Multiracial candidates run in districts with comparable White and Latino VAP shares (White: 51%/53%, Latino: 8%/7%),¹⁵ there is evidence that they run in districts with smaller Black VAP shares (15%/24%) than their Monoracial peers. This gap is most apparent in majority-Black districts, at 6% to 19%. Yet Multiracial candidates run in majority-minority districts at rates equal to their Monoracial peers (50%/47%). This suggests that while identifying as Multiracial does not preclude running in a majority-minority district, it tracks a meaningfully different constituency profile. The pattern is consistent with Multiracial identification signaling reach into more racially heterogeneous coalitions rather than concentrated Black electorates.

15. Throughout, paired figures report Multiracial/Monoracial values, in that order.




Figure 5: Multiracial share of Black-identified candidates by state. *Note:* Every state is an equal-size tile; fill is the Multiracial share of Black candidate-cycles (grayscale, light = low), and each tile prints the state code and its candidate-cycle count (n) so thin-support states are visible.

Table 5: District Composition by Identity Group: Median / [IQR] and District-Type Shares

Measure	Multiracial	Monoracial
White VAP share (%)	49.5 [43.4, 69.5]	52.5 [36.3, 70.0]
Black VAP share (%)	15.2 [6.8, 36.3]	24.0 [13.5, 46.2]
Latino VAP share (%)	8.8 [4.7, 18.5]	6.4 [3.6, 15.9]
Dem. pres. vote share (%)	63.6 [49.5, 77.0]	61.8 [44.2, 75.8]
Majority-Black district	5.9%	18.7%
Majority-minority district	51.5%	46.6%

Note: Voting-age-population (VAP) measures and Democratic presidential vote share are median / [25th, 75th percentile]; majority-Black and majority-minority district are group shares.

While the district composition and geographical variation in where Multiracial candidates run is distinct compared to Monoracials, a critical pathway to winning is campaign funding (Magee 2002). Significantly for our argument, if the framing of Multiracial candidates as strategically engaging with their identity to build multiethnic coalitions holds (Lemi 2017, 2026), we should observe its imprint in campaign financing as well. Table 6 shows the median campaign contributions and the percentages receiving institutional endorsements, separated into Black-specific gatekeepers (CBC PAC, Collective PAC, and Higher Heights) and general institutional funders (the DCCC and EMILY’s List). While descriptive, the evidence points to substantial variation. On median total receipts, individual contributions, PAC contributions, and number of distinct givers, Multiracial candidates outpace Monoracial candidates. In short, the descriptive gradient is consistent with a financial premium attached to Multiracial identification.

The sharper pattern lies in endorsements from Black-specific gatekeepers. While Multiracial candidates draw more from donors broadly, the endorsement gap is widest among the Black-specific gatekeepers. Some 29% received a CBC PAC endorsement versus roughly 14% of Monoracials, with a similar (though smaller) edge for Collective PAC and a nearly 15-point gap in receiving any Black-specific gatekeeper endorsement (CBC or Collective). General institutional funders show smaller, though nonzero, gaps, a comparison Table 6 makes directly. This concentration is what distinguishes a general electability story from a strategic-signaling one, since the endorsement premium is largest precisely where Black-specific gatekeepers decide whom to back. Consistent with this, Multiracials win their primaries about 55% of the time versus 49% for their Monoracial peers, though this raw gap does not yet net out the district and resource differences documented above. By identifying as Multiracial, candidates appear to signal viability to institutional insiders; in backing them, Black-specific gatekeepers assume the risk that a candidate may lack the racial competency necessary to substantively represent Black interests (Johnson 2026).

We argue that institutional gatekeeping benefits Multiracial candidates broadly, as they are framed as being able to build cross-racial coalitions and thus win their primaries at higher rates. However, because gender shapes both candidate entry and institutional gatekeeping (Crowder-Meyer 2013; Dittmar, Sanbonmatsu, and Carroll 2018; Sanbonmatsu 2002), we hypothesized that the Multiracial premium would be gendered, accruing primarily to men. Table 7 provides an overview of campaign funding and institutional endorsement by gender. Multiracial men have a clear contribution advantage. Across number of donors, individual contributions, and PAC contributions, Multiracial male candidates outpace every other group. Interestingly, we see that the same cannot be said for Monoracial men. In fact, Monoracial women have a larger donor network, higher individual donations, and larger PAC donations than their male counterparts. Multiracial women present a mixed profile, with a larger donor count than Monoracial Black men but weaker individual and PAC contributions. Taken together, the contribution data support the core prediction of our intersectional gatekeeping argument: the Multiracial premium accrues to men.

When examining institutional endorsements, a similar pattern arises. Again, Multiracial men are endorsed by the CBC PAC at far higher rates than Monoracial men (36%/11%), over

Table 6: Pre-Primary Funding and Institutional Support by Multiracial Identification: Descriptive Comparison (LLM Coding)

Outcome	Multiracial (n=68)	Monoracial (n=920)	Gap
<i>Pre-primary fundraising</i>			
Pre-Primary Receipts	\$904,625 [\$27,717, \$1,947,745]	\$114,183 [\$5,411, \$1,068,696]	\$790,442
Pre-Primary Individual Contribs.	\$279,398 [\$25,786, \$1,144,694]	\$69,338 [\$4,796, \$353,222]	\$210,060
Pre-Primary PAC Contribs.	\$167,554 [\$436, \$1,137,834]	\$18,225 [\$0, \$582,077]	\$149,329
Pre-Primary Distinct Givers	1,508 [70, 6,459]	293 [20, 1,674]	1,214
<i>Black-specific institutional support</i>			
CBC PAC	29.4%	14.0%	+15.4 pp
Collective PAC	17.6%	9.0%	+8.6 pp
Any Black gatekeeper (CBC/Collective)	32.4%	17.2%	+15.2 pp
Higher Heights	11.8%	6.7%	+5.0 pp
<i>General institutional support</i>			
DCCC	19.1%	10.4%	+8.7 pp
EMILY's List	7.4%	5.8%	+1.6 pp
<i>Primary outcome</i>			
Won Primary	55.2%	48.8%	+6.4 pp

Note: Dollar and donor outcomes are group medians with [25th, 75th percentile]; institutional support and primary success are group shares. Gap is the raw Multiracial-minus-Monoracial difference (medians for dollars/counts; percentage points for shares). Black-specific institutional support: CBC PAC, Collective PAC, their union, and Higher Heights (Black women); general institutional support: DCCC and EMILY's List (institutional but not Black-specific). Dollar and donor outcomes are pre-primary (received on or before the FEC nomination date).

triple the share of candidates receiving a CBC PAC endorsement. However, both Monoracial and Multiracial women also maintain larger shares of candidates with CBC PAC endorsements than Monoracial men, at 17% and 19%, respectively, compared to 11%. Finally, Multiracial men dominate in primary victories, with roughly 67% of candidates winning their election. A similar pattern arises when examining the three other groups: Monoracial women again come in second in terms of primary victories (52%), followed by Monoracial men (47%) and, finally, Multiracial women (33%).

The gender results also raise a puzzle. The Multiracial premium accrues to men, as our intersectional gatekeeping argument predicts. Yet women candidates, Monoracial women in particular, do better in most categories than Monoracial men, an advantage our theory alone does not anticipate. However, our data allow us to examine the role of the #MeToo movement as well as the historic rise in Black women's candidacies during the 2018 midterm election season (Simien 2022, 2023). If these events constituted unique shocks to the funding and institutional endorsement of Black women, their effects should be visible when we examine

Table 7: Pre-Primary Funding and Institutional Support by Identity and Gender: Descriptive Comparison (LLM Coding)

Measure	Multi. men (n=44)	Mono. men (n=545)	Multi. women (n=24)	Mono. women (n=375)
<i>Pre-primary fundraising</i>				
Total receipts (\$)	\$1,074,759 [\$144,136, \$3,330,080]	\$81,296 [\$1,678, \$929,989]	\$65,060 [\$6,219, \$1,331,056]	\$164,299 [\$12,192, \$1,212,903]
Individual contribs. (\$)	\$425,847 [\$75,427, \$1,431,999]	\$58,630 [\$1,450, \$287,009]	\$45,022 [\$5,921, \$696,333]	\$93,926 [\$9,407, \$432,477]
PAC contribs. (\$)	\$545,185 [\$18,154, \$1,289,063]	\$13,400 [\$0, \$532,261]	\$5,900 [\$0, \$330,902]	\$31,268 [\$1,125, \$693,772]
Donors (#)	1,580 [612, 7,198]	227 [9, 1,284]	352 [17, 4,816]	471 [41, 2,226]
<i>Black-specific institutional support</i>				
CBC PAC support (%)	36.4%	10.8%	16.7%	18.7%
Collective PAC support (%)	15.9%	5.3%	20.8%	14.4%
Any Black gatekeeper (%)	36.4%	12.5%	25.0%	24.0%
Higher Heights (%)	0.0%	0.4%	33.3%	16.0%
<i>General institutional support</i>				
DCCC (%)	25.0%	9.5%	8.3%	11.7%
EMILY's List (%)	6.8%	4.0%	8.3%	8.3%
<i>Primary outcome</i>				
Won primary (%)	67.4%	46.9%	33.3%	51.6%

Note: Purely descriptive. Donor and dollar measures are median / [IQR] and are pre-primary (on or before the FEC nomination date). Institutional rows are the candidate-file SUPPORT flags: the Black-specific gatekeepers (CBC PAC, Collective PAC, their union, and Higher Heights) and the general institutional funders (DCCC and EMILY's List) – the same measures as the institutional panels of Table 6. Any Black gatekeeper is the CBC/Collective union. Cell counts appear in the column headers. No hypothesis tests are reported.

the change in endorsement and funding strategies from 2016 through 2024.¹⁶

Figure 6 shows the endorsement rate and the median contributions/donors by year and institution type (Black-specific versus general). Endorsement trends (top panel) show a dramatic increase in the percentage of women receiving Black-specific support from the CBC (6%), Collective (16%), and Higher Heights (19%) PACs. (See Appendix Table A2 for the numerical percentage-point changes.) In contrast, men either outright lose endorsements (CBC: -6%) or gain at lower rates than women over the same period.¹⁷ For general institutional endorsements, there is a less stark change over the period, consistent with the historic 2018 period inducing a race-specific shock within Black-specific endorsements. A similar trend holds across all funding and donor sources (lower panel).¹⁸ These findings together provide evidence that the relative decline in male funding and endorsements was a product of the #MeToo movement and historic gains by Black women in the lead-up to 2018.

16. While the Collective PAC was founded in 2016, all other PACs predate our data collection. Therefore, this is not a story about PACs being founded explicitly to fund women in our shock period. Rather, we are able to descriptively show how the shock altered institutional funding and endorsement strategies.

17. Multiracial men's high endorsement levels (Table 7) are not inconsistent with this downward trend; the crossover reflects change over the period rather than absolute rates of support.

18. Primary win rates also decline more steeply for men than for women during this period. While women also lose their primaries, the difference between male and female candidates highlights how gains in endorsements and funding track greater primary success for women.




Figure 6: Gendered realignment of institutional support. *Note:* Endorsement and primary-success rates by funder (top row) and median pre-primary funding and donor outcomes (bottom row), by gender and cycle. The post-2016 crossover is concentrated among the Black-specific gatekeepers (CBC PAC, Collective PAC, Higher Heights); general institutional funders (DCCC, EMILY's List) show parallel movement across genders.

6 Conclusion

Multiracial identity in electoral behavior is not uniformly beneficial. Our findings suggest that gender plays a significant role in undermining the idea that Multiracial candidates are uniquely positioned to build coalitions and win elections (Lemi 2017, 2021, 2026; Velasquez-Manoff 2017). While Multiracial men win at 67%, Multiracial women secure their party’s nomination just 33% of the time. As a result, archetypal depictions of Obama as a bridge across racial divisions were overgeneralized and overstated. For Multiracial men, racial and gender signals interact in a mutually enforcing and beneficial way; yet, these same signals produce a different evaluative process that ultimately leads to defeat for Multiracial women.

Part of the reason Multiracial women face hurdles is due to Institutional gatekeepers. Institutions, in particular, PACs and other donors, interpret candidates through a gendered reality (Crowder-Meyer 2013; Dittmar, Sanbonmatsu, and Carroll 2018; Sanbonmatsu 2006). Thus, we build a theory of intersectional gatekeeping that argues political institutions selectively engage with candidates who are likely to win elections (Crenshaw 1989). Historically, this has been men (Crowder-Meyer 2013; Sanbonmatsu 2002), and results indicate that PACs and campaign contributions still reflect this idea. PAC expenditures, endorsements, and contributions broadly defined are highest for Multiracial men, while Multiracial women receive no equivalent premium. Without this intersectional argument, the general takeaway would be that Multiracial candidates are homogeneous in their ability to win nominations; however, using our unique dataset, we find that this is not the case.

Significantly, the very institutions producing the intersectional gatekeeping are the same Black-specific political organizations that reinvested in women beginning in the 2018 midterm elections. The consolidation around Black women by the CBC PAC, Collective PAC, and Higher Heights is explained by the #MeToo movement as well as the symbolic rise of African American female candidates (Simien 2022, 2023; Smooth 2021). This investment in Black women, however, circumscribed Multiracial female candidates, as they were least likely to receive Black organizational endorsements and funding. As a result, the narrowing of internal diversity among Black women representatives has strong implications for descriptive representation (Dovi 2002; Mansbridge 1999; Wamble 2025).

Without the creation of an LLM-based classification system that expanded data collection on racial identity to Multiracials, discovering these nuances would be limited to the handful of mixed-race incumbents (Hardy-Fanta et al. 2013; Lemi 2017). However, while LLM performance can exceed human-based coding (Benoit et al. 2025; Gilardi, Alizadeh, and Kubli 2023; Heseltine and Clemm von Hohenberg 2024; Törnberg 2023), we agree that LLM output should be measured against a gold standard (Grimmer and Stewart 2013). Validation techniques show that our API-based coding scheme largely matches our RA-based racial coding scheme. The skill file, thus, is a useful measurement tool for classifying political candidates’ racial identity (Halterman and Keith 2026).

Future research should expand the validation process of LLM-assisted racial identity coding to include a broader spectrum of models and techniques. In particular, varying

the temperature of the model, whether the LLM is proprietary or open-source, and the language utilized in the skill file can all affect the intended outcome (Benoit et al. 2025; McLaren et al. 2026; Ollion et al. 2024; Pangakis, Wolken, and Fasching 2023; Spirling 2023). Researchers seeking to fully test our coding methodology should expand the validation process to account for these variations in modeling. Moreover, future research will have to grapple with Multiracial and Monoracial identities beyond the Black-White paradigm. Expanding the classification system to include a larger set of racial categories is paramount to research in representation, political economy, and the electoral behavior of federal-level elites broadly (Branton 2009; Fraga 2016; Fraga and Hassell 2021; Grumbach and Sahn 2020).¹⁹ Finally, our method would benefit from comparisons between LLM-based models and traditional machine learning techniques.

As Multiracial candidates continue to proliferate within the political arena, researchers will need to grapple with proper ways to identify them. Moreover, findings here indicate that Multiracial identification will become increasingly strategic—deployed to secure endorsements and funding, and to cultivate perceptions of candidates as coalition builders. Yet identity is intersectional. Without examining Multiracial and Monoracial Black candidates through this lens, we limit our ability to fully analyze the scope conditions under which identification is strategically beneficial or a folly. Only by embracing new measurement tools can we begin to unpack these nuanced questions.

19. Expanding our dataset to state-level office is also a fruitful opportunity for future research. While the number of candidates at the federal level are low, state and local elections are where a majority of Multiracial candidates are currently located (Hardy-Fanta et al. 2013; Lemi 2017).

References

- Argyle, Lisa P, and Michael Barber. 2024. "Misclassification and bias in predictions of individual ethnicity from administrative records." *American Political Science Review* 118 (2): 1058–1066.
- Atsusaka, Yuki. 2021. "A logical model for Predicting Minority representation: Application to redistricting and Voting Rights cases." *American Political Science Review* 115 (4): 1210–1225.
- Benoit, Kenneth, Scott De Marchi, Conor Laver, Michael Laver, and Jinshuai Ma. 2025. "Using large language models to analyze political texts through natural language understanding." *American Journal of Political Science*.
- Bonica, Adam. 2017. "Professional networks, early fundraising, and electoral success." *Election Law Journal: Rules, Politics, and Policy* 16 (1): 153–171.
- Branton, Regina P. 2009. "The Importance of Race and Ethnicity in Congressional Primary Elections." *Political Research Quarterly* 62 (3): 459–473.
- Brown, Nadia. 2014a. "'It's more than hair... that's why you should care': the politics of appearance for Black women state legislators." *Politics, Groups, and Identities* 2 (3): 295–312.
- Brown, Nadia E. 2014b. *Sisters in the statehouse: Black women and legislative decision making*. Oxford University Press.
- Brown, Nadia E, Christopher J Clark, and Anna Mahoney. 2022. "The Black women of the US congress: learning from descriptive data." *Journal of Women, Politics & Policy* 43 (3): 328–346.
- Clopper, C. J., and E. S. Pearson. 1934. "The Use of Confidence or Fiducial Limits Illustrated in the Case of the Binomial." *Biometrika* 26 (4): 404–413.
- Cohen, Jacob. 1960. "A Coefficient of Agreement for Nominal Scales." *Educational and Psychological Measurement* 20 (1): 37–46.
- Collins, Patricia Hill. 2000. "Gender, black feminism, and black political economy." *The annals of the American academy of political and social science* 568 (1): 41–53.
- Crenshaw, Kimberle. 1989. "Demarginalizing the intersection of race and gender in antidiscrimination law, feminist theory, and antiracist politics." 140:139–167.
- Crowder-Meyer, Melody. 2013. "Gendered recruitment without trying: How local party recruiters affect women's representation." *Politics & Gender* 9 (4): 390–413.
- Dasanaik, Noah. 2026. "Large Language Models Naively Recover Ethnicity from Individual Records." *arXiv preprint arXiv:2601.21132*.
- Davenport, Lauren D. 2016. "Beyond Black and White: Biracial Attitudes in Contemporary U.S. Politics." *American Political Science Review*, <https://doi.org/10.1017/s0003055415000556>.

- Davis, F James. 1991. *Who is black?: One nation's definition*. Penn State Press.
- Dawson, Michael C. 1994. *Behind the mule: Race and class in African-American politics*. Princeton University Press.
- De Winter, Joost CF, Samuel D Gosling, and Jeff Potter. 2016. "Comparing the Pearson and Spearman correlation coefficients across distributions and sample sizes: A tutorial using simulations and empirical data." *Psychological methods* 21 (3): 273.
- Denton, Nancy A, and Glenn D Deane. 2010. "Researching race and ethnicity: Methodological issues." *The SAGE handbook of race and ethnic studies*. London: Sage, 67–89.
- Dittmar, Kelly, Kira Sanbonmatsu, and Susan J Carroll. 2018. *A seat at the table: Congresswomen's perspectives on why their presence matters*. Oxford University Press.
- Dovi, Suzanne. 2002. "Preferable descriptive representatives: Will just any woman, black, or Latino do?" *American Political Science Review* 96 (4): 729–743.
- Feinstein, Alvan R., and Domenic V. Cicchetti. 1990. "High Agreement but Low Kappa: I. The Problems of Two Paradoxes." *Journal of Clinical Epidemiology* 43 (6): 543–549.
- Fields, Karen E., and Barbara J. Fields. 2012. *Racecraft: The Soul of Inequality in American Life*. London: Verso.
- Fraga, Bernard L. 2016. "Candidates or districts? Reevaluating the role of race in voter turnout." *American Journal of Political Science* 60 (1): 97–122.
- Fraga, Bernard L, Eric Gonzalez Juenke, and Paru Shah. 2020. "One run leads to another: Minority incumbents and the emergence of lower ticket minority candidates." *The Journal of Politics* 82 (2): 771–775.
- Fraga, Bernard L, and Hans JG Hassell. 2021. "Are minority and women candidates penalized by party politics? Race, gender, and access to party support." *Political Research Quarterly* 74 (3): 540–555.
- Fraga, Bernard L, Paru Shah, and Eric Gonzalez Juenke. 2020. "Did women and candidates of color lead or ride the democratic wave in 2018?" *PS: Political Science & Politics* 53 (3): 435–439.
- Gilardi, Fabrizio, Meysam Alizadeh, and Maël Kubli. 2023. "ChatGPT outperforms crowd workers for text-annotation tasks." *Proceedings of the National Academy of Sciences* 120 (30): e2305016120.
- Goggin, Stephen N. 2017. "Partisan yet Personal Politicians: Candidates' Biographies and Why They Matter." *Book Manuscript*.
- Greengard, Philip, and Andrew Gelman. 2026. "A calibrated BISG for inferring race from surname and geolocation." *Journal of the Royal Statistical Society Series A: Statistics in Society* 189 (1): 512–543.

- Grimmer, Justin, and Brandon M Stewart. 2013. "Text as data: The promise and pitfalls of automatic content analysis methods for political texts." *Political analysis* 21 (3): 267–297.
- Grumbach, Jacob M, and Alexander Sahn. 2020. "Race and Representation in Campaign Finance." *American Political Science Review* 114 (1): 206–221.
- Gwet, Kilem L. 2008. "Computing Inter-Rater Reliability and Its Variance in the Presence of High Agreement." *British Journal of Mathematical and Statistical Psychology* 61 (1): 29–48.
- Halterman, Andrew, and Katherine A Keith. 2026. "Codebook llms: Evaluating llms as measurement tools for political science concepts." *Political Analysis* 34 (2): 188–204.
- Hardy-Fanta, Carol, Pei-te Lien, Dianne Marie Pinderhughes, and Christine M Sierra. 2013. "Racial and Ethnic Identity of Elected Officials of Color: A Closer Look at a Complex Matter." In *APSA 2013 Annual Meeting Paper, American Political Science Association 2013 Annual Meeting*.
- Hassell, Hans JG, and Neil Visalvanich. 2019. "The Party's Primary preferences: race, gender, and party support of Congressional Primary Candidates." *American Journal of Political Science* 63 (4): 905–919.
- Hauke, Jan, and Tomasz Kossowski. 2011. "Comparison of values of Pearson's and Spearman's correlation coefficients on the same sets of data." *Quaestiones geographicae* 30 (2): 87–93.
- Heseltine, Michael, and Bernhard Clemm von Hohenberg. 2024. "Large language models as a substitute for human experts in annotating political text." *Research & Politics* 11 (1): 20531680241236239.
- Imai, Kosuke, and Kabir Khanna. 2016. "Improving ecological inference by predicting individual ethnicity from voter registration records." *Political Analysis* 24 (2): 263–272.
- Imai, Kosuke, Santiago Olivella, and Evan TR Rosenman. 2022. "Addressing census data problems in race imputation via fully Bayesian Improved Surname Geocoding and name supplements." *Science Advances* 8 (49): eadc9824.
- Johnson, Isaiah C. 2026. "Jack of All Trades, Master of None: Multiracial Politicians and Political Representation."
- Juenke, Eric Gonzalez. 2014. "Ignorance is Bias: The Effect of Latino Losers on Models of Latino Representation." *American Journal of Political Science* 58 (3): 593–603.
- Juenke, Eric Gonzalez, and Paru Shah. 2016. "Demand and Supply: Racial and Ethnic Minority Candidates in White Districts." *Journal of Race, Ethnicity, and Politics* 1 (1): 60–90.
- Lemi, Danielle Casarez. 2017. "Identity and coalitions in a multiracial era: how state legislators navigate race and ethnicity." *Politics, Groups, and Identities* 6:725–742. ISSN: 2156-5503, 2156-5511. <https://dx.doi.org/10.1080/21565503.2017.1288144>.

- Lemi, Danielle Casarez. 2021. “Do voters prefer just any descriptive representative? The case of multiracial candidates.” *Perspectives on Politics* 19 (4): 1061–1081.
- . 2026. *Doing Identity Labor: How Mixed-Race Politicians Disrupt Descriptive Representation*. Oxford University Press.
- Lemi, Danielle Casarez, and Nadia E. Brown. 2019. “Melanin and Curls: Evaluation of Black Women Candidates.” *Journal of Race, Ethnicity, and Politics* 4 (2): 259–296.
- Leslie, Gregory, and David O Sears. 2022. “The heaviest drop of blood: Black exceptionalism among multiracials.” *Political Psychology* 43 (6): 1123–1145.
- Leslie, Gregory John. 2022. *Colored by Context: Relative Racial Salience (RRS) and the Politics of Mixed-Race Americans*. University of California, Los Angeles.
- Lockhart, Jeffrey W, Molly M King, and Christin Munsch. 2023. “Name-based demographic inference and the unequal distribution of misrecognition.” *Nature Human Behaviour* 7 (7): 1084–1095.
- Lublin, David, and Matthew Wright. 2024. “Diversity Matters: The Election of Asian Americans to U.S. State and Federal Legislatures.” *American Political Science Review* 118 (1): 380–400.
- Magee, Christopher. 2002. “Do political action committees give money to candidates for electoral or influence motives?” *Public Choice* 112 (3): 373–399.
- Mansbridge, Jane. 1999. “Should blacks represent blacks and women represent women? A contingent” yes.” *The Journal of politics* 61 (3): 628–657.
- McLaren, Lorcan, James Cross, Zuzanna Krakowska, Robin Rauner, and Martijn Schoonvelde. 2026. “Magic Words or Methodical Work? Challenging Conventional Wisdom in LLM-Based Political Text Annotation.” *arXiv preprint arXiv:2603.26898*.
- Meisels, Mellissa. 2026. “Everything in moderation? The effect of extremist nominations on individual and corporate PAC fundraising.” *Political Science Research and Methods* 14 (2): 353–369.
- Myers, Andrew CW. 2025. “Do Donors Punish Extremist Primary Nominees? Evidence from Congress and American State Legislatures.” *American Political Science Review* 1:21.
- Ollion, Étienne, Rubing Shen, Ana Macanovic, and Arnault Chatelain. 2024. “The dangers of using proprietary LLMs for research.” *Nature Machine Intelligence* 6 (1): 4–5.
- Omi, Michael, and Howard Winant. 1986. *Racial Formation in the United States: From the 1960s to the 1980s*. New York: Routledge & Kegan Paul.
- Pangakis, Nicholas, Samuel Wolken, and Neil Fasching. 2023. “Automated annotation with generative ai requires validation.” *arXiv preprint arXiv:2306.00176*.
- Philpot, Tasha S, and Hanes Walton Jr. 2007. “One of our own: Black female candidates and the voters who support them.” *American Journal of Political Science* 51 (1): 49–62.

- Porter, Rachel, and Sarah A Treul. 2025. "Evaluating (in) experience in congressional elections." *American Journal of Political Science* 69 (1): 284–298.
- Reingold, Beth, Kerry L Haynie, and Kirsten Widner. 2020. *Race, gender, and political representation: Toward a more intersectional approach*. Oxford University Press.
- Rendleman, Hunter E. 2023. "Bound Together: Racial Peer Effects and Caucus Control in the US Congress."
- Rockquemore, Kerry Ann, and David L. Brunnsma. 2008. *Beyond Black: Biracial Identity in America*. 2nd ed. Lanham, MD: Rowman & Littlefield Publishers.
- Sanbonmatsu, Kira. 2002. "Gender stereotypes and vote choice." *american Journal of political Science*, 20–34.
- . 2006. "Do parties know that "women win"? Party leader beliefs about women's electoral chances." *Politics & Gender* 2 (4): 431–450.
- Sen, Maya, and Omar Wasow. 2016. "Race as a bundle of sticks: Designs that estimate effects of seemingly immutable characteristics." *Annual Review of Political Science* 19 (1): 499–522.
- Shah, Paru. 2014. "It takes a black candidate: A supply-side theory of minority representation." *Political Research Quarterly* 67 (2): 266–279.
- Shah, Paru, Jamil Scott, and Eric Gonzalez Juenke. 2019. "Women of Color Candidates: Examining Emergence and Success in State Legislative Elections." *Politics, Groups, and Identities* 7 (2): 429–443.
- Simien, Evelyn M. 2022. *Historic firsts in US elections: trailblazing candidates in gubernatorial, congressional, and mayoral campaigns*. Routledge.
- . 2023. "The 2018 Congressional Midterms, Symbolic Empowerment, and Ayanna Pressley's Mobilizing Effect: A Case Study for Future Analysis of Historic Firsts." In *Women of Color Political Elites in the US*, 17–34. Routledge.
- Slaughter, Christine, Chaya Crowder, and Christina Greer. 2024. "Black women: Keepers of democracy, the democratic process, and the Democratic Party." *Politics & Gender* 20 (1): 162–181.
- Smooth, Wendy. 2021. "Splitting Condi (licious)." *Re-Imagining Black Women: A Critique of Post-Feminist and Post-Racial Melodrama in Culture and Politics*, 27.
- Snipp, C Matthew. 2003. "Racial measurement in the American census: Past practices and implications for the future." *Annual Review of Sociology* 29 (1): 563–588.
- Spencer, Rainier. 2004. "Assessing multiracial identity theory and politics: The challenge of hypodescent." *Ethnicities* 4 (3): 357–379.
- Spirling, Arthur. 2023. "Why open-source generative AI models are an ethical way forward for science." *Nature* 616 (7957): 413–413.

- Thomas, Anita Jones, and Suzette L Speight. 1999. “Racial identity and racial socialization attitudes of African American parents.” *Journal of Black psychology* 25 (2): 152–170.
- Thomsen, Danielle M. 2023. “Competition in congressional elections: Money versus votes.” *American Political Science Review* 117 (2): 675–691.
- Törnberg, Petter. 2023. “Chatgpt-4 outperforms experts and crowd workers in annotating political twitter messages with zero-shot learning.” *arXiv preprint arXiv:2304.06588*.
- Velasquez-Manoff, Moisés. 2017. “What Biracial People Know.” *The New York Times*, https://www.nytimes.com/2017/03/04/opinion/sunday/what-biracial-people-know.html?_r=0.
- Wamble, Julian J. 2025. *We Choose You: How Black Voters Decide Which Candidates to Support*. Cambridge University Press.
- White, Ismail K, and Chryl N Laird. 2020. “Steadfast democrats: How social forces shape Black political behavior.” In *Steadfast Democrats*. Princeton University Press.
- Ziems, Caleb, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. “Can large language models transform computational social science?” *Computational Linguistics* 50 (1): 237–291.

Appendix Contents

A Additional Descriptive Results	34
B LLM Validation: Robustness	36
B.1 Concurrent validity	36
B.2 Predictive validity	36
B.3 Gold-standard replication	36

A Additional Descriptive Results

Table A1: Ideology and CF Score by Identity Group: Median / [IQR]

Measure	Multiracial	Monoracial
CF score (recipient)	-0.75 [-0.97, 0.79]	-0.71 [-0.99, -0.13]
CF score (dynamic)	-0.75 [-0.94, 0.33]	-0.73 [-1.00, -0.12]
DW-NOMINATE (1st dim.)	-0.42 [-0.45, -0.33]	-0.45 [-0.49, -0.36]

Note: Recipient CF score, dynamic CF score, and first-dimension DW-NOMINATE by identity group, median / [25th, 75th percentile]. DW-NOMINATE exists only for candidates who served and is sparser.

Table A2: Gendered Realignment of Institutional Endorsements: 2016 Versus 2018–2024

Funder / group	Gender	2016	2018–2024	Change
<i>Pre-primary fundraising (medians)</i>				
Total receipts (\$)	Men	\$68,456	\$110,311	\$41,856
Total receipts (\$)	Women	\$61,728	\$164,244	\$102,516
Individual contribs. (\$)	Men	47,392	67,934	20,543
Individual contribs. (\$)	Women	41,640	103,849	62,209
PAC contribs. (\$)	Men	12,700	16,678	3,978
PAC contribs. (\$)	Women	12,556	28,256	15,700
Donors (#)	Men	103	307	204
Donors (#)	Women	107	514	408
<i>Black-specific institutional support</i>				
CBC PAC	Men	18.1%	11.9%	-6.2 pp
CBC PAC	Women	12.8%	19.2%	+6.3 pp
Collective PAC	Men	1.2%	6.9%	+5.7 pp
Collective PAC	Women	0.0%	16.4%	+16.4 pp
Higher Heights	Men	0.0%	0.4%	+0.4 pp
Higher Heights	Women	0.0%	18.9%	+18.9 pp
<i>General institutional support</i>				
DCCC	Men	6.0%	11.5%	+5.4 pp
DCCC	Women	7.7%	11.9%	+4.3 pp
EMILY's List	Men	1.2%	4.7%	+3.5 pp
EMILY's List	Women	5.1%	8.6%	+3.5 pp
<i>Primary outcome</i>				
Won primary	Men	61.4%	46.2%	-15.2 pp
Won primary	Women	56.4%	49.9%	-6.6 pp
<i>Any Black gatekeeper, by identity-gender group</i>				
Any Black gatekeeper: Monoracial men		14.3% (n=77)	12.2% (n=468)	-2.1 pp
Any Black gatekeeper: Monoracial women		13.5% (n=37)	25.1% (n=338)	+11.6 pp
Any Black gatekeeper: Multiracial men		66.7% (n=6)	31.6% (n=38)	-35.1 pp
Any Black gatekeeper: Multiracial women		0.0% (n=2)	27.3% (n=22)	+27.3 pp

Note: Top panel: median pre-primary funding and donor outcomes by gender and era. Middle panels: share of candidates endorsed by each funder, by gender and era. Bottom panel: any Black gatekeeper (CBC PAC or Collective PAC) endorsement rate by identity-gender group; 2016 Multiracial cells are very small and should be read with caution. The 2016/2018 boundary coincides with the #MeToo movement, the 2018 surge in Black women's candidacies, and Higher Heights' PAC scale-up.

B LLM Validation: Robustness

B.1 Concurrent validity

Table A3: Concurrent Validity of the LLM Race Coding Against the RA Gold Standard

Statistic	Count	Estimate	95% CI
Precision (PPV)	55/68	0.809	[0.695, 0.894]
Recall (sensitivity)	55/57	0.965	[0.879, 0.996]
Specificity	918/931	0.986	[0.976, 0.993]
Accuracy	973/988	0.985	[0.975, 0.991]
F1 (Multiracial class)		0.880	[0.813, 0.935]
Cohen’s κ		0.872	[0.803, 0.930]
Gwet’s AC1		0.983	[0.973, 0.991]
Exact McNemar (FP vs. FN)	13 vs. 2	$p = 0.007^{***}$	

Note: The gold standard is the research-assistant coding, complete for all candidate-cycles; metrics are computed on the head-to-head sample of LLM-coded candidate-cycles. Exact (Clopper–Pearson) CIs for the proportion-based classification metrics; percentile-bootstrap CIs ($B = 2000$) over candidate-cycles for F1, κ , and AC1, which are not binomial proportions (κ and AC1 share one set of resamples). AC1 is the prevalence-robust companion to κ , reported because the Multiracial base rate ($\approx 6\%$) places κ in the range where the kappa paradox is a known concern. The exact McNemar test is a two-sided binomial test on the discordant pairs; the excess of false positives over false negatives indicates the LLM errs toward over-inclusion, so LLM-coded gaps are attenuated relative to the gold standard. * $p < 0.1$; ** $p < 0.05$; *** $p < 0.01$

B.2 Predictive validity

Table A4: Predictive Validity of the LLM Coding: Downstream-Gap Consistency

Statistic	Level	Estimate	95% CI
Downstream-gap consistency: Pearson r	Outcome battery (11 outcomes)	0.925	[0.430, 0.996]
Downstream-gap consistency: Spearman ρ	Outcome battery (11 outcomes)	0.973	[0.445, 0.991]

Note: Rows correlate standardized Multiracial-minus-Monoracial gaps across the outcome battery under the two codings (continuous outcomes: median gap scaled by the pooled IQR; binary outcomes: percentage-point share gaps); see Figure 3. Both CIs are percentile bootstraps over candidate-cycles ($B = 2000$), recomputing the full gap vectors within each resample; analytic (Fisher- z) intervals over the 11 outcomes would be invalid because the gaps are dependent statistics of the same sample rather than independent observations.

B.3 Gold-standard replication

Table A5: Pre-Primary Funding and Institutional Support: Descriptive Comparison Under the RA Gold Standard (Same Sample)

Outcome	Multiracial (n=57)	Monoracial (n=931)	Gap
<i>Pre-primary fundraising</i>			
Pre-Primary Receipts	\$1,189,829 [\$101,250, \$2,769,197]	\$109,980 [\$5,477, \$1,046,583]	\$1,079,849
Pre-Primary Individual Contribs.	\$437,502 [\$57,379, \$1,280,886]	\$66,721 [\$4,869, \$351,826]	\$370,781
Pre-Primary PAC Contribs.	\$521,562 [\$3,800, \$1,317,941]	\$17,650 [\$0, \$563,711]	\$503,912
Pre-Primary Distinct Givers	1,988 [389, 6,768]	289 [20, 1,650]	1,699
<i>Black-specific institutional support</i>			
CBC PAC	33.3%	14.0%	+19.4 pp
Collective PAC	19.3%	9.0%	+10.3 pp
Any Black gatekeeper (CBC/Collective)	35.1%	17.2%	+17.9 pp
Higher Heights	14.0%	6.7%	+7.4 pp
<i>General institutional support</i>			
DCCC	21.1%	10.4%	+10.6 pp
EMILY's List	10.5%	5.6%	+4.9 pp
<i>Primary outcome</i>			
Won Primary	62.5%	48.4%	+14.1 pp

Note: Dollar and donor outcomes are group medians with [25th, 75th percentile]; institutional support and primary success are group shares. Gap is the raw Multiracial-minus-Monoracial difference (medians for dollars/counts; percentage points for shares). Black-specific institutional support: CBC PAC, Collective PAC, their union, and Higher Heights (Black women); general institutional support: DCCC and EMILY's List (institutional but not Black-specific). All institutional measures are candidate-file support flags. Dollar and donor outcomes are pre-primary (received on or before the FEC nomination date).

Table A6: Sample Overview Under the RA Gold Standard (Same Sample)

Measure	Multiracial	Monoracial	Total
All candidate-cycles	57	931	988
Share of coded sample	5.8%	94.2%	100.0%
Unique candidates			540
Men	35	554	589
Women	22	377	399
% Democrat	73.7%	75.6%	

Note: Counts of candidate-cycles and unique candidates, the group shares of the coded sample, the gender split, and the share identifying as Democrats, under the RA gold-standard coding on the head-to-head sample.

Table A7: District Composition by Identity Group Under the RA Gold Standard

Measure	Multiracial	Monoracial
White VAP share (%)	47.6 [39.2, 68.4]	52.9 [36.3, 70.2]
Black VAP share (%)	23.0 [8.5, 44.7]	23.8 [13.3, 46.2]
Latino VAP share (%)	9.1 [5.8, 22.1]	6.4 [3.6, 15.9]
Dem. pres. vote share (%)	63.6 [51.3, 77.2]	61.8 [44.2, 75.8]
Majority-Black district	10.5%	18.3%
Majority-minority district	57.9%	46.2%

Note: Voting-age-population (VAP) measures and Democratic presidential vote share are median / [25th, 75th percentile]; majority-Black and majority-minority district are group shares.

Table A8: Pre-Primary Funding and Institutional Support by Identity and Gender Under the RA Gold Standard

Measure	Multi. men (n=35)	Mono. men (n=554)	Multi. women (n=22)	Mono. women (n=377)
<i>Pre-primary fundraising</i>				
Total receipts (\$)	\$1,528,945 [\$904,625, \$4,111,200]	\$75,664 [\$1,834, \$917,895]	\$64,779 [\$1,292, \$1,340,378]	\$164,299 [\$12,343, \$1,207,378]
Individual contribs. (\$)	\$678,697 [\$144,205, \$2,325,134]	\$56,388 [\$1,459, \$286,949]	\$43,819 [\$1,181, \$864,845]	\$93,926 [\$10,300, \$430,101]
PAC contribs. (\$)	\$886,241 [\$202,514, \$1,453,604]	\$12,836 [\$0, \$518,345]	\$5,900 [\$0, \$336,150]	\$31,268 [\$1,000, \$684,965]
Donors (#)	2,320 [796, 17,788]	224 [9, 1,286]	613 [5, 5,374]	460 [43, 2,206]
<i>Black-specific institutional support</i>				
CBC PAC support (%)	42.9%	10.8%	18.2%	18.6%
Collective PAC support (%)	20.0%	5.2%	18.2%	14.6%
Any Black gatekeeper (%)	42.9%	12.5%	22.7%	24.1%
Higher Heights (%)	0.0%	0.4%	36.4%	15.9%
<i>General institutional support</i>				
DCCC (%)	28.6%	9.6%	9.1%	11.7%
EMILY's List (%)	8.6%	4.0%	13.6%	8.0%
<i>Primary outcome</i>				
Won primary (%)	82.4%	46.3%	31.8%	51.6%

Note: Donor and dollar measures are median / [IQR] and are pre-primary (on or before the FEC nomination date). Institutional rows are the candidate-file SUPPORT flags: the Black-specific gatekeepers (CBC PAC, Collective PAC, their union, and Higher Heights) and the general institutional funders (DCCC and EMILY's List). Any Black gatekeeper is the CBC/Collective union. Cell counts appear in the column headers.

Table A9: Gendered Realignment of Institutional Endorsements Under the RA Gold Standard: 2016 Versus 2018–2024

Funder / group	Gender	2016	2018–2024	Change
<i>Pre-primary fundraising (medians)</i>				
Total receipts (\$)	Men	\$68,456	\$110,311	\$41,856
Total receipts (\$)	Women	\$61,728	\$164,244	\$102,516
Individual contribs. (\$)	Men	47,392	67,934	20,543
Individual contribs. (\$)	Women	41,640	103,849	62,209
PAC contribs. (\$)	Men	12,700	16,678	3,978
PAC contribs. (\$)	Women	12,556	28,256	15,700
Donors (#)	Men	103	307	204
Donors (#)	Women	107	514	408
<i>Black-specific institutional support</i>				
CBC PAC	Men	18.1%	11.9%	-6.2 pp
CBC PAC	Women	12.8%	19.2%	+6.3 pp
Collective PAC	Men	1.2%	6.9%	+5.7 pp
Collective PAC	Women	0.0%	16.4%	+16.4 pp
Higher Heights	Men	0.0%	0.4%	+0.4 pp
Higher Heights	Women	0.0%	18.9%	+18.9 pp
<i>General institutional support</i>				
DCCC	Men	6.0%	11.5%	+5.4 pp
DCCC	Women	7.7%	11.9%	+4.3 pp
EMILY's List	Men	1.2%	4.7%	+3.5 pp
EMILY's List	Women	5.1%	8.6%	+3.5 pp
<i>Primary outcome</i>				
Won primary	Men	61.4%	46.2%	-15.2 pp
Won primary	Women	56.4%	49.9%	-6.6 pp
<i>Any Black gatekeeper, by identity-gender group</i>				
Any Black gatekeeper: Monoracial men		14.3% (n=77)	12.2% (n=477)	-2.1 pp
Any Black gatekeeper: Monoracial women		13.9% (n=36)	25.2% (n=341)	+11.3 pp
Any Black gatekeeper: Multiracial men		66.7% (n=6)	37.9% (n=29)	-28.7 pp
Any Black gatekeeper: Multiracial women		0.0% (n=3)	26.3% (n=19)	+26.3 pp

Note: Top panel: median pre-primary funding and donor outcomes by gender and era. Middle panels: share of candidates endorsed by each funder, by gender and era. Bottom panel: any Black gatekeeper (CBC PAC or Collective PAC) endorsement rate by identity-gender group; 2016 Multiracial cells are very small and should be read with caution. The 2016/2018 boundary coincides with the #MeToo movement, the 2018 surge in Black women's candidacies, and Higher Heights' PAC scale-up.

Table A10: Ideology and CF Score by Identity Group Under the RA Gold Standard

Measure	Multiracial	Monoracial
CF score (recipient)	-0.66 [-0.95, -0.25]	-0.71 [-1.00, -0.13]
CF score (dynamic)	-0.71 [-0.92, -0.18]	-0.73 [-1.00, -0.11]
DW-NOMINATE (1st dim.)	-0.42 [-0.45, -0.33]	-0.45 [-0.49, -0.36]

Note: Recipient CF score, dynamic CF score, and first-dimension DW-NOMINATE by identity group, median / [25th, 75th percentile]. DW-NOMINATE exists only for candidates who served and is sparser.